

```
library(class)
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
df <- read.csv("UniversalBank.csv") #Assigning data to the variable
head(df)
```

```
##   ID Age Experience Income ZIP.Code Family CCAvg Education Mortgage
## 1  1  25          1    49    91107      4   1.6          1          0
## 2  2  45         19    34    90089      3   1.5          1          0
## 3  3  39         15    11    94720      1   1.0          1          0
## 4  4  35          9   100    94112      1   2.7          2          0
## 5  5  35          8    45    91330      4   1.0          2          0
## 6  6  37         13    29    92121      4   0.4          2        155
##   Personal.Loan Securities.Account CD.Account Online CreditCard
## 1              0                  1          0      0          0
## 2              0                  1          0      0          0
## 3              0                  0          0      0          0
## 4              0                  0          0      0          0
## 5              0                  0          0      0          1
## 6              0                  0          0      1          0
```

```
dim(df)
```

```
## [1] 5000  14
```

```
t(t(names(df))) # This function creates a transpose of the dataframe
```

```
##      [,1]
## [1,] "ID"
## [2,] "Age"
## [3,] "Experience"
## [4,] "Income"
## [5,] "ZIP.Code"
## [6,] "Family"
## [7,] "CCAvg"
## [8,] "Education"
## [9,] "Mortgage"
## [10,] "Personal.Loan"
## [11,] "Securities.Account"
## [12,] "CD.Account"
## [13,] "Online"
## [14,] "CreditCard"
```

Dropping ID and ZIP Code columns

```
df <- df[,-c(1,5)]
head(df)
```

```
##   Age Experience Income Family CCAvg Education Mortgage Personal.Loan
## 1  25          1     49      4   1.6          1          0           0
## 2  45         19     34      3   1.5          1          0           0
## 3  39         15     11      1   1.0          1          0           0
## 4  35          9    100      1   2.7          2          0           0
## 5  35          8     45      4   1.0          2          0           0
## 6  37         13     29      4   0.4          2        155           0
##   Securities.Account CD.Account Online CreditCard
## 1                   1           0         0          0
## 2                   1           0         0          0
## 3                   0           0         0          0
## 4                   0           0         0          0
## 5                   0           0         0          1
## 6                   0           0         1          0
```

Splitting Data to 60% training and 40% validation. Before we split, let us transform categorical variables into dummy variables

```
# Education converted into factor
df$Education <- as.factor(df$Education)

# converting Education into Dummy Variables
groups <- dummyVars(~., data = df) # This creates the dummy groups
df_m <- as.data.frame(predict(groups,df))
```

```
#Split Data into 60% training and 40% validation
set.seed(1)
train.index <- sample(row.names(df_m), 0.6*dim(df_m)[1])
valid.index <- setdiff(row.names(df_m), train.index)
train.df <- df_m[train.index,]
valid.df <- df_m[valid.index,]
```

Normalizing train data & validation data

```
train.norm.df <- train.df[,-10] #Personal Loan is the 10th variable
valid.norm.df <- valid.df[,-10]

norm.values <- preProcess(train.df[, -10], method=c("center", "scale"))
train.norm.df <- predict(norm.values, train.df[, -10])
valid.norm.df <- predict(norm.values, valid.df[, -10])
```

#Questions

Consider the following customer: 1. Age = 40, Experience = 10, Income = 84, Family = 2, CCAvg = 2, Education\_1 = 0, Education\_2 = 1, Education\_3 = 0, Mortgage = 0, Securities Account = 0, CD Account = 0, Online = 1, and Credit Card = 1. Perform a k-NN classification with all predictors except ID and ZIP code using k = 1. Remember to transform categorical predictors with more than two categories into dummy variables first. Specify the success class as 1 (loan acceptance), and use the default cutoff value of 0.5. How would this customer be classified?

```

# creating a new sample
new_customer <- data.frame(
  Age = 40,
  Experience = 10,
  Income = 84,
  Family = 2,
  CCAvg = 2,
  Education.1 = 0,
  Education.2 = 1,
  Education.3 = 0,
  Mortgage = 0,
  Securities.Account = 0,
  CD.Account = 0,
  Online = 1,
  CreditCard = 1
)
# Normalize the new customer
new_cust.norm <- new_customer
new_cust.norm <- predict(norm.values, new_cust.norm)

```

```

knn.pred1 <- class::knn(train = train.norm.df,
  test = new_cust.norm, cl = train.df$Personal.Loan, k = 1, prob = TRUE)
knn.pred1

```

```

## [1] 0
## attr(,"prob")
## [1] 1
## Levels: 0 1

```

The new\_customer data where k=1 the model is classified as 0, Hence the new customer does not take the personal loan.

2. What is a choice of k that balances between overfitting and ignoring the predictor information?

```

# Calculating the accuracy for each value of k
# Setting the range of k values to consider
accuracy.df <- data.frame(k = seq(1, 15, 1), overallaccuracy = rep(0, 15))
for(i in 1:15) {
  knn.pred <- class::knn(train = train.norm.df,
    test = valid.norm.df, cl = train.df$Personal.Loan, k = i)

  accuracy.df[i, 2] <- confusionMatrix(knn.pred, as.factor(valid.df$Personal.Loan), positive = "1")$overall
}

accuracy.df

```

```

##      k overallaccuracy
## 1    1          0.9630
## 2    2          0.9565
## 3    3          0.9640
## 4    4          0.9595
## 5    5          0.9605

```

```
## 6 6 0.9575
## 7 7 0.9580
## 8 8 0.9575
## 9 9 0.9535
## 10 10 0.9550
## 11 11 0.9495
## 12 12 0.9475
## 13 13 0.9500
## 14 14 0.9500
## 15 15 0.9485
```

```
which(accuracy.df[,2] == max(accuracy.df[,2]))
```

```
## [1] 3
```

Best k value is 3 that balances between overfitting and ignoring the predictor information.

3. Show the confusion matrix for the validation data that results from using the best k.

```
prediction_knn <- class::knn(train = train.norm.df, test = valid.norm.df, cl = train.df$Personal.Loan, k
confusionMatrix(prediction_knn, as.factor(valid.df$Personal.Loan))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 1786   63
##           1    9  142
##
##           Accuracy : 0.964
##           95% CI : (0.9549, 0.9717)
##           No Information Rate : 0.8975
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.7785
##
##           McNemar's Test P-Value : 4.208e-10
##
##           Sensitivity : 0.9950
##           Specificity : 0.6927
##           Pos Pred Value : 0.9659
##           Neg Pred Value : 0.9404
##           Prevalence : 0.8975
##           Detection Rate : 0.8930
##           Detection Prevalence : 0.9245
##           Balanced Accuracy : 0.8438
##
##           'Positive' Class : 0
##
```

The confusion matrix for validation set using k=3 results: Accuracy : 0.964, Sensitivity : 0.9950, Specificity : 0.6927

4. Consider the following customer: Age = 40, Experience = 10, Income = 84, Family = 2, CCAvg = 2, Education\_1 = 0, Education\_2 = 1, Education\_3 = 0, Mortgage = 0, Securities Account = 0, CD Account = 0, Online = 1 and Credit Card = 1. Classify the customer using the best k.

```
# creating a new sample
new_customer <- data.frame(
  Age = 40,
  Experience = 10,
  Income = 84,
  Family = 2,
  CCAvg = 2,
  Education.1 = 0,
  Education.2 = 1,
  Education.3 = 0,
  Mortgage = 0,
  Securities.Account = 0,
  CD.Account = 0,
  Online = 1,
  CreditCard = 1
)
# Normalize the new customer
new.cust.norm <- new_customer
new.cust.norm <- predict(norm.values, new.cust.norm)

knn.pred1 <- class::knn(train = train.norm.df,
  test = new.cust.norm,
  cl = train.df$Personal.Loan, k = 3)
knn.pred1
```

```
## [1] 0
## Levels: 0 1
```

The new\_customer data where k=3 the model is classified as 0, Hence the new customer does not take the personal loan.

5. Repartition the data, this time into training, validation, and test sets (50% : 30% : 20%). Apply the k-NN method with the k chosen above. Compare the confusion matrix of the test set with that of the training and validation sets. Comment on the differences and their reason

```
# Repartition the training, validation and test sets to 50,30, and 20 percents.
set.seed(1)
train.index = sample(row.names(df_m), 0.5*dim(df_m)[1])
remaining.index = setdiff(row.names(df_m), train.index)
valid.index = sample(remaining.index, 0.3*dim(df_m)[1])
test.index = setdiff(remaining.index, valid.index)

#Loading the partitioned data into the dataframe.
train.df = df_m[train.index,]
valid.df= df_m[valid.index,]
test.df = df_m[test.index,]

#Normalizing the data after repartitioning accordingly.
train.norm.df <- train.df[, -10]
```

```

valid.norm.df <- valid.df[, -10]
test.norm.df <- test.df[, -10]

norm.values <- preProcess(train.df[, -10], method = c("center", "scale"))
train.norm.df <- predict(norm.values, train.df[, -10])
valid.norm.df <- predict(norm.values, valid.df[, -10])
test.norm.df <- predict(norm.values, test.df[, -10])

#Applying the k-NN method to all the sets that we have using k value as 3.

#Training set
knn_train <- class::knn(train = train.norm.df, test = train.norm.df, cl = train.df$Personal.Loan, k = 3)

#Validation set
knn_valid <- class::knn(train = train.norm.df, test = valid.norm.df, cl = train.df$Personal.Loan, k = 3)

#Test set
knn_test <- class::knn(train = train.norm.df, test = test.norm.df, cl = train.df$Personal.Loan, k = 3)

#Confusion Matrix

#Training set confusion matrix
confusionMatrix(knn_train, as.factor(train.df$Personal.Loan))

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 2263   54
##           1    5  178
##
##           Accuracy : 0.9764
##           95% CI : (0.9697, 0.982)
##       No Information Rate : 0.9072
##       P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8452
##
##  McNemar's Test P-Value : 4.129e-10
##
##           Sensitivity : 0.9978
##           Specificity : 0.7672
##       Pos Pred Value : 0.9767
##       Neg Pred Value : 0.9727
##           Prevalence : 0.9072
##       Detection Rate : 0.9052
##       Detection Prevalence : 0.9268
##       Balanced Accuracy : 0.8825
##
##       'Positive' Class : 0
##

```

```
#Validation set confusion matrix
confusionMatrix(knn_valid, as.factor(valid.df$Personal.Loan))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 1358   42
##           1    6   94
##
##           Accuracy : 0.968
##           95% CI : (0.9578, 0.9763)
##       No Information Rate : 0.9093
##       P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.7797
##
##  McNemar's Test P-Value : 4.376e-07
##
##           Sensitivity : 0.9956
##           Specificity : 0.6912
##       Pos Pred Value : 0.9700
##       Neg Pred Value : 0.9400
##           Prevalence : 0.9093
##       Detection Rate : 0.9053
##   Detection Prevalence : 0.9333
##       Balanced Accuracy : 0.8434
##
##       'Positive' Class : 0
##
```

```
#Testing set confusion matrix
confusionMatrix(knn_test, as.factor(test.df$Personal.Loan))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0  884   35
##           1    4   77
##
##           Accuracy : 0.961
##           95% CI : (0.9471, 0.9721)
##       No Information Rate : 0.888
##       P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.777
##
##  McNemar's Test P-Value : 1.556e-06
##
##           Sensitivity : 0.9955
##           Specificity : 0.6875
##       Pos Pred Value : 0.9619
```

```

##          Neg Pred Value : 0.9506
##          Prevalence : 0.8880
##          Detection Rate : 0.8840
##          Detection Prevalence : 0.9190
##          Balanced Accuracy : 0.8415
##
##          'Positive' Class : 0
##

```

Miscalculation:  $M = \text{false positive} + \text{false negative}$

Training set=59, Validation set=48, Testing set=39

Accuracy: Training set=0.9764, Validation set=0.968, Testing set=0.961 The accuracy on the training set is slightly higher than on the validation and test sets, but all sets have high accuracy, indicating that the model performs well in terms of overall correctness.

Sensitivity: Training set=0.9978, Validation set=0.9956, Testing set=0.9955 Sensitivity measures the models ability to detect the positive class (in this case, class 1) properly. All of the sets have extremely high sensitivity, suggesting that the model is exceptionally good at spotting class 1 cases.

Specificity: Training set=0.7672, Validation set=0.6912, Testing set=0.6875 Specificity measures the models ability to accurately identify the negative class (class 0 in this example). The training set has the maximum specificity, but the test and validation sets have lower specificity values, implying that the model is less accurate at recognizing class 0 occurrences.

Training set confusion matrix results: Miscalculations: 59, Accuracy : 0.9764, Sensitivity : 0.9978, Specificity : 0.7672

Validation set confusion matrix results: Miscalculations: 48, Accuracy : 0.968, Sensitivity : 0.9956, Specificity : 0.6912

Testing set confusion matrix results : Miscalculations:39, Accuracy : 0.961, Sensitivity : 0.9955, Specificity : 0.6875

In summary, the model exhibits good overall performance across the training, validation, and test sets, with slight variations in metrics. Notably, there is a noticeable decrease in specificity as we move from the training set to the validation and test sets. This drop suggests that the model may have a higher tendency to make false positive predictions (classifying instances as class 1 when they are actually class 0) on unseen data.